

#### Implementation:

1. Pick a random number from [0,7) which will decide how we jitter the input
2. 0 means normal, the implementation that was present in blue - turn the 784 inputs into a 784x1 np.array
3. 1 means shift right: Go through lists of 28 in the dataset (one row), at each row, np.roll it over by 1. Add the rolled version to a final list that is later converted into an np.array and reshape into a 784x1 input array.
4. 2 means to shift left. Do the same as Step 3 however np.roll it over by -1.
5. 3 means to shift up. This time when we go through the lists of 28, we are not adjusting them, we are simply adding them to a final list that is a 28x28 list of lists. Turn that into an np.array and then np.roll it by -1 on axis = 0 (vertical shifting)
6. 4 means to shift down. Do the same as Step 5 however np.roll it by 1 still on the same axis
7. 5 means rotate to the right 15 degrees. Very similar to Step 5 in creating the list of lists structure. Then use `scipy.ndimage.rotate()` with `reshape=False` (very important because this ensures that we don't change the size of the dataset) to rotate the array 15 degrees and then convert it into an np.array and normalize it.
8. 6 means to shift to the left 15 degrees. Do the same as Step 7 but `scipy.ndimage.rotate()` by -15 degrees.

#### Comparing:

Whoa. While the code was running, I was almost certain that my outcome from “jittering” would be worse. As seen in the screenshot below, 30 epochs started off very similarly in both networks (distorted and undistorted), but at about the 8th or 9th epoch diverged greatly. Undistorted continued to decrease at an almost constant rate, and got as low as 0.97% inaccuracy on the training set as it was training. The distorted network? While it decreased, it was MUCH slower. It got down to about 2.81% inaccuracy pretty rapidly and in a similar pace, but after that by the end of 30 epochs only got to 2.45% inaccuracy on the training set while it trained. However, this is not where the story ends. I then ran both networks on the TEST set. Here is where I was shocked. The undistorted network got 3.21% inaccuracy, which is pretty good, but wildly higher than the inaccuracy “predicted” while training (should have been sub 1%). As for the distorted network, well that got a 2.5% inaccuracy! Not a drastic improvement, but way closer to the “predicted” accuracy (2.45%) and still LOWER than the undistorted set. Cool stuff!

Note: Network architecture was identical for both cases - [784, 300, 100, 10]

0 6.836666666666667  
1 5.293333333333333  
2 3.705  
3 3.385  
4 3.236666666666667  
5 3.141666666666667  
6 3.051666666666667  
7 2.961666666666664  
8 2.903333333333333  
9 2.851666666666666  
10 2.813333333333335  
11 2.773333333333334  
12 2.756666666666667  
13 2.73  
14 2.696666666666668  
15 2.663333333333333  
16 2.631666666666664  
17 2.608333333333334  
18 2.588333333333334  
19 2.568333333333334  
20 2.541666666666665  
21 2.531666666666667  
22 2.506666666666667  
23 2.510000000000002  
24 2.503333333333334  
25 2.493333333333333  
26 2.483333333333334  
27 2.469999999999998  
28 2.469999999999998  
29 2.461666666666664  
30 2.455